

# Configuring Network Policies for Applications

Network Policies allow you to limit connections between Pods. Therefore, using Network Policies provide better security by reducing the compromise radius.

**Note** Network Policies determine whether a connection is allowed, and they do not offer higher level features like authorization or secure transport (like SSL/TLS).

## Create a cluster

To create a container cluster with Network Policy enforcement, run the following command:

```
gcloud container clusters create admatic-network-test --enable-network-policy
```

**Warning** If you omit the `--enable-network-policy` flag, any `NetworkPolicy` resources you create are silently ignored.

WARNING: Starting in 1.12, new clusters will have basic authentication disabled by default. Basic authentication can be enabled (or disabled) manually using the `--[no-]enable-basic-auth` flag.

WARNING: Starting in 1.12, new clusters will not have a client certificate issued. You can manually enable (or disable) the issuance of the client certificate using the `--[no-]issue-client-certificate` flag.

WARNING: Currently VPC-native is not the default mode during cluster creation. In the future, this will become the default mode and can be disabled using `--no-enable-ip-alias` flag. Use `--[no-]enable-ip-alias` flag to suppress this warning.

WARNING: Starting in 1.12, default node pools in new clusters will have their legacy Compute Engine instance metadata endpoints disabled by default. To create a cluster with legacy instance metadata endpoints disabled in the default node pool, run `clusters create` with the flag `--metadata disable-legacy-endpoints=true`.

This will enable the autorepair feature for nodes. Please see <https://cloud.google.com/kubernetes-engine/docs/node-auto-repair> for more information on node autorepairs.

WARNING: Starting in Kubernetes v1.10, new clusters will no longer get compute-rw and storage-ro scopes added to what is specified in `--scopes` (though the latter will remain included in the default `--scopes`). To use these scopes, add them explicitly to `--scopes`. To use the new behavior, set `container/new_scopes_behavior` property (`gcloud config set container/new_scopes_behavior true`).

Creating cluster admatic-network-test in us-west1-c...done.

Created [<https://container.googleapis.com/v1/projects/espblufi-android/zones/us-west1-c/clusters/admatic-network-test>].

To inspect **the** contents **of** your cluster, go **to**: [https://console.cloud.google.com/kubernetes/workload\\_/gcloud/us-west1-c/admatic-network-test?project=espblufi-android](https://console.cloud.google.com/kubernetes/workload_/gcloud/us-west1-c/admatic-network-test?project=espblufi-android) kubeconfig entry generated **for** admatic-network-test.

NAME	LOCATION	MASTER_VERSION	MASTER_IP	MACHINE_TYPE	NOD
E_VERSION	NUM_NODES	STATUS			
admatic-network-test	us-west1-c	1.10.9-gke.5	35.233.184.239	n1-standard-1	1.1
0.9-gke.5	3	RUNNING			

## Restrict incoming traffic to Pods

Kubernetes NetworkPolicy resources let you configure network access policies for the Pods. NetworkPolicy objects contains with the following information:

- Pods to which the network policies apply, usually designated by a label selector
- Type of Internet traffic the network policy affects: Ingress for incoming traffic, Egress for outgoing traffic, or both
- For Ingress policies, which Pods can connect to the specified Pods
- For Egress policies, the Pods to which the specified Pods can connect.

First, run a simple web server application with label `app=hello` and expose it internally in the cluster:

```
kubectl run hello-web --labels app=hello \
  --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose
```

```
kubectl run --generator=deployment/apps.v1 is DEPRECATED and will be removed in a fu
ture version. Use kubectl run --generator=run-pod/v1 or kubectl create instead.
service/hello-web created
deployment.apps/hello-web created
```

Next, you need to configure a NetworkPolicy to allow traffic to the `hello-web` Pods from only the `app=foo` Pods. Other incoming traffic from Pods that do not have this label, external traffic, and traffic from Pods in other namespaces are blocked.

```
cat << EOF > hello-allow-from-foo.yaml
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: hello-allow-from-foo
spec:
  policyTypes:
    - Ingress
  podSelector:
    matchLabels:
      app: hello
  ingress:
```

```
- from:
  - podSelector:
      matchLabels:
        app: foo
EOF
```

This policy selects Pods with label `app=hello` and specifies an ingress policy to allow traffic only from Pods with the label `app=foo`.

To apply this policy to the cluster, run the following command:

```
kubectl apply -f hello-allow-from-foo.yaml
```

```
networkpolicy.networking.k8s.io/hello-allow-from-foo created
```

## Validate the ingress policy

First, run a temporary Pod with the label `app=foo` and get a shell in the Pod:

```
kubectl run -l app=foo --image=alpine --restart=Never --rm -i -t test-1
```

*If you don't see a command prompt, try pressing enter.*  
/ #

Make a request to the `hello-web:8080` endpoint to verify that the incoming traffic is allowed:

```
wget -qO- --timeout=2 http://hello-web:8080
```

```
Hello, world!
Version: 1.0.0
Hostname: hello-web-8b44b849-9c2v1
```

Traffic from Pod `app=foo` to the `app=hello` Pods is enabled.

Next, run a temporary Pod with a different label (`app=other`) and get a shell inside the Pod:

```
kubectl run -l app=other --image=alpine --restart=Never --rm -i -t test-1
```

*If you don't see a command prompt, try pressing enter.*  
/ #

Make the same request to observe that the traffic is **not allowed** and therefore the request times out, then exit from the Pod shell:

```
wget -q0- --timeout=2 http://hello-web:8080
```

```
wget: download timed out
```

```
exit
```

```
pod "test-1" deleted
pod default/test-1 terminated (Error)
```

## Restrict outgoing traffic from the Pods

You can restrict outgoing (egress) traffic as you would incoming traffic.

However, to be able to query internal hostnames such as `hello-web` or external hostnames such as `www.example.com`, you must allow DNS (domain name system) resolution in your egress network policies. DNS traffic occurs on port 53 using TCP and UDP protocols.

To exercise egress network policies, deploy a `NetworkPolicy` controlling outbound traffic from Pods with the label `app=foo` while allowing traffic only to Pods with the label `app=hello`, as well as the DNS traffic.

```
cat << EOF > foo-allow-to-hello.yaml
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: foo-allow-to-hello
spec:
  policyTypes:
    - Egress
  podSelector:
    matchLabels:
      app: foo
  egress:
    - to:
        - podSelector:
            matchLabels:
              app: hello
    - ports:
        - port: 53
          protocol: TCP
        - port: 53
          protocol: UDP
EOF

kubectl apply -f foo-allow-to-hello.yaml
```

```
networkpolicy.networking.k8s.io/foo-allow-to-hello created
```

This manifest specifies a network policy controlling the egress traffic from Pods with label `app=foo` with two allowed destinations:

1. Pods in the same namespace with the label `app=hello`.
2. Cluster Pods or external endpoints on port 53 (UDP and TCP)

## Validate the egress policy

First, deploy a new web application called `hello-web-2` and expose it internally in the cluster:

```
kubectl run hello-web-2 --labels app=hello-2 \
  --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose
```

```
kubectl run --generator=deployment/apps.v1 is DEPRECATED and will be removed in a fu
ture version. Use kubectl run --generator=run-pod/v1 or kubectl create instead.
service/hello-web-2 created
deployment.apps/hello-web-2 created
```

Next, run a temporary Pod with `app=foo` label and get a shell prompt inside the container:

```
kubectl run -l app=foo --image=alpine --rm -i -t --restart=Never test-3
```

*If you don't see a command prompt, try pressing enter.*  
/ #

Validate that the Pod can establish connections to `hello-web:8080`:

```
wget -qO- --timeout=2 http://hello-web:8080
```

```
Hello, world!
Version: 1.0.0
Hostname: hello-web-8b44b849-9c2v1
```

Validate that the Pod cannot establish connections to `hello-web-2:8080`:

```
wget -qO- --timeout=2 http://hello-web-2:8080
```

```
wget: download timed out
```

```
exit
```

```
pod "test-3" deleted
pod default/test-3 terminated (Error)
```

# Cleaning up

## Delete the container cluster

This step will delete the resources that make up the container cluster, such as the compute instances, disks and network resources.

```
gcloud container clusters delete admatic-network-test
```

The following clusters will be deleted.

- [admatic-network-**test**] in [us-west1-c]

Do you want to **continue** (Y/n)? **Y**

Deleting **cluster** admatic-network-**test**...done.

Deleted [<https://container.googleapis.com/v1/projects/espblufi-android/zones/us-west1-c/clusters/admatic-network-test>].