

ReplicaSet

Initially, ReplicationControllers were the only Kubernetes component for replicating pods and rescheduling them when nodes failed. Later, a similar resource called a ReplicaSet was introduced. It's a new generation of ReplicationController and replaces it completely (ReplicationControllers will eventually be deprecated).

You usually won't create them directly, but instead have them created automatically when you create the higher-level Deployment resource.

ReplicaSet is the next-generation Replication Controller. The only difference between a ReplicaSet and a Replication Controller right now is the selector support. ReplicaSet supports the new set-based selector requirements as described in the labels user guide whereas a Replication Controller only supports equality-based selector requirements.

Comparing a ReplicaSet to a ReplicationController

A ReplicaSet behaves exactly like a ReplicationController, but it has more expressive pod selectors. Whereas a ReplicationController's label selector only allows matching pods that include a certain label, a ReplicaSet's selector also allows matching pods that lack a certain label or pods that include a certain label key, regardless of its value.

Also, for example, a single ReplicationController can't match pods with the label `env=production` and those with the label `env=devel` at the same time. It can only match either pods with the `env=production` label or pods with the `env=devel` label. But a single ReplicaSet can match both sets of pods and treat them as a single group.

Similarly, a ReplicationController can't match pods based merely on the presence of a label key, regardless of its value, whereas a ReplicaSet can. For example, a ReplicaSet can match all pods that include a label with the key `env`, whatever its actual value is (you can think of it as `env=*`).

Defining a ReplicaSet

```
cat << EOF > kubia-replicaset.yaml
apiVersion: apps/v1beta2
kind: ReplicaSet
metadata:
  name: kubia
spec:
```

```
replicas: 3
selector:
  matchLabels:
    app: kubia
template:
  metadata:
    labels:
      app: kubia
  spec:
    containers:
      - name: kubia
        image: luksa/kubia
EOF
```

The first thing to note is that ReplicaSets aren't part of the v1 API, so you need to ensure you specify the proper `apiVersion` when creating the resource.

You're creating a resource of type `ReplicaSet` which has much the same contents as the `Replication-Controller`. The only difference is in the selector. Instead of listing labels the pods need to have directly under the selector property, you're specifying them under `selector.matchLabels`. This is the simpler (and less expressive) way of defining label selectors in a `ReplicaSet`. Later, you'll look at the more expressive option, as well.

About the API version attribute

This is your first opportunity to see that the `apiVersion` property specifies two things:

- The API group (which is `apps` in this case)
- The actual API version (`v1beta2`)

Certain Kubernetes resources are in what's called the core API group, which doesn't need to be specified in the `apiVersion` field (you just specify the version—for example, you've been using `apiVersion: v1` when defining Pod resources). Other resources, which were introduced in later Kubernetes versions, are categorized into several API groups.

Creating and examining a ReplicaSet

Create the `ReplicaSet` from the YAML file with the `kubectl create` command. After that, you can examine the `ReplicaSet` with `kubectl get` and `kubectl describe`:

```
kubectl create -f kubia-replicaset.yaml
```

```
replicaset.apps "kubia" created
```

Use `rs` shorthand, which stands for `replicaset`.

```
kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubia	3	3	3	2m

```
kubectl describe rs
```

```
Name:          kubia
Namespace:     default
Selector:      app=kubia
Labels:        app=kubia
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=kubia
  Containers:
    kubia:
      Image:      luksa/kubia
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:      <none>
      Volumes:      <none>
Events:
  Type     Reason             Age   From                    Message
  ----     -
  Normal   SuccessfulCreate   3m    replicaset-controller   Created pod: kubia-8n5f4
  Normal   SuccessfulCreate   3m    replicaset-controller   Created pod: kubia-rf9fq
  Normal   SuccessfulCreate   3m    replicaset-controller   Created pod: kubia-xgccl
```

As you can see, the ReplicaSet isn't any different from a ReplicationController. It's showing it has three replicas matching the selector.

```
kubectl delete rs kubia
```

```
replicaset.extensions "kubia" deleted
```

Deleting the ReplicaSet should delete all the pods.

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubia-8n5f4	1/1	Terminating	0	8m
kubia-rf9fq	1/1	Terminating	0	8m
kubia-xgccl	1/1	Terminating	0	8m

Using the ReplicaSet's more expressive label selectors

The main improvements of ReplicaSets over ReplicationControllers are their more expressive label selectors. You intentionally used the simpler `matchLabels` selector in the first ReplicaSet example to see that ReplicaSets are no different from Replication-Controllers. Now, you'll rewrite the selector to use the more powerful `matchExpressions` property, as shown in the following listing.

```
cat << EOF > kubia-replicaset-matchexpressions.yaml
apiVersion: apps/v1beta2
kind: ReplicaSet
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    matchExpressions:
      - key: app
        operator: In
        values:
          - kubia
  template:
    metadata:
      labels:
        app: kubia
    spec:
      containers:
        - name: kubia
          image: luksa/kubia
EOF
```

You can add additional expressions to the selector. As in the example, each expression must contain a key, an operator, and possibly (depending on the operator) a list of values. You'll see four valid operators:

- `In`—Label's value must match one of the specified values.
- `NotIn`—Label's value must not match any of the specified values.
- `Exists`—Pod must include a label with the specified key (the value isn't important). When using this operator, you shouldn't specify the values field.
- `DoesNotExist`—Pod must not include a label with the specified key. The values property must not be specified.

If you specify multiple expressions, all those expressions must evaluate to true for the selector to match a pod. If you specify both `matchLabels` and `matchExpressions`, all the labels must match and all the expressions must evaluate to true for the pod to match the selector.

```
kubectl create -f kubia-replicaset-matchexpressions.yaml
```

```
replicaset.apps "kubia" created
```

```
kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubia	3	3	3	18s

```
kubectl describe rs
```

```
Name:          kubia
Namespace:     default
Selector:      app in (kubia)
Labels:        app=kubia
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=kubia
  Containers:
    kubia:
      Image:      luksa/kubia
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
  Volumes:      <none>
Events:
  Type    Reason             Age    From                      Message
  ----    -
  Normal  SuccessfulCreate   49s    replicaset-controller     Created pod: kubia-z56zt
  Normal  SuccessfulCreate   49s    replicaset-controller     Created pod: kubia-k9nxp
  Normal  SuccessfulCreate   49s    replicaset-controller     Created pod: kubia-5jnc2
```

```
kubectl delete rs kubia
```

```
replicaset.extensions "kubia" deleted
```